

PERFORMANCE IMPROVEMENTS IN SQL SERVER 2019

Ohio North SQL Server User Group
December 1, 2020

ABOUT ME

In IT for 21 years, 13 of it working with SQL Server
Senior Data Engineer at Concurrency, Inc.

Learning more about Azure every day

Lover of all things internal to SQL Server

When not working with the Microsoft Data Platform I love to read, volunteer at the Art Institute of Chicago, and hang out with my cat

Twitter - @skreebydba

Email - skreebydba@gmail.com

Blog - www.skreebydba.com



WHAT WE WILL COVER

Transaction log

Current crash recovery process

Accelerated Database Recovery (ADR)

ADR crash recovery process

In-memory tempdb metadata

Persistent Memory



TRANSACTIONS

Unit of work in the database

All transactions begin

Transactions can commit or rollback

Default behavior is auto-commit



TRANSACTION LOG

Records all changes to the database

Changes written to the log buffer in memory

SQL Server uses write-ahead logging (WAL)

Log buffer is flushed to disk on COMMIT or when it fills up

Data pages associated with transactions can remain in memory



TRANSACTION LOG ARCHITECTURE

Transaction log contains logical units called Virtual Log Files (VLF)

VLFs can be active or free

VLFs containing log records that may be needed must be active

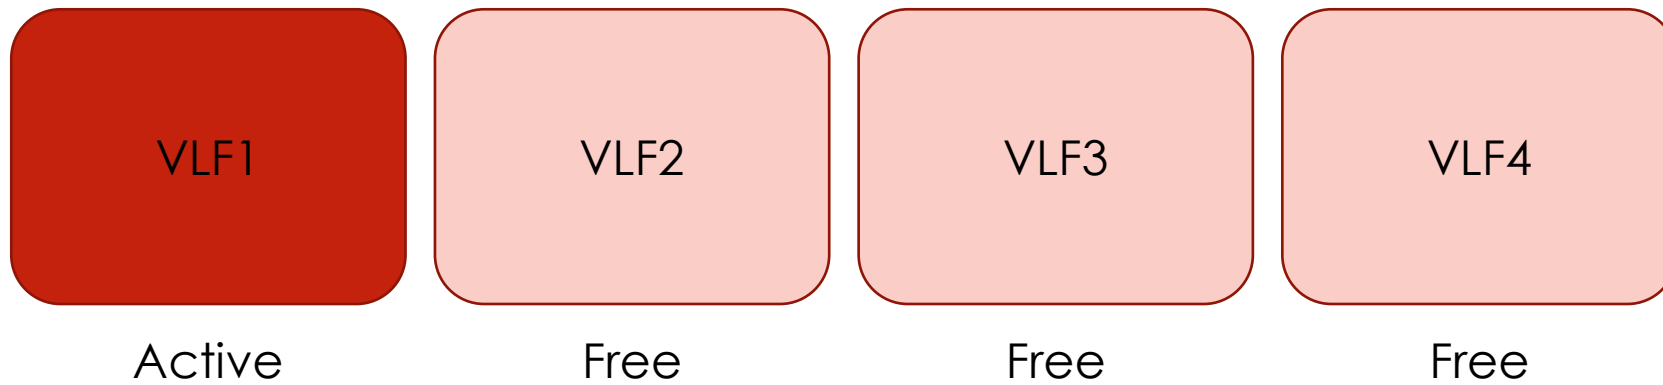
Always one active VLF

In SIMPLE recovery, VLFs freed by CHECKPOINT operation

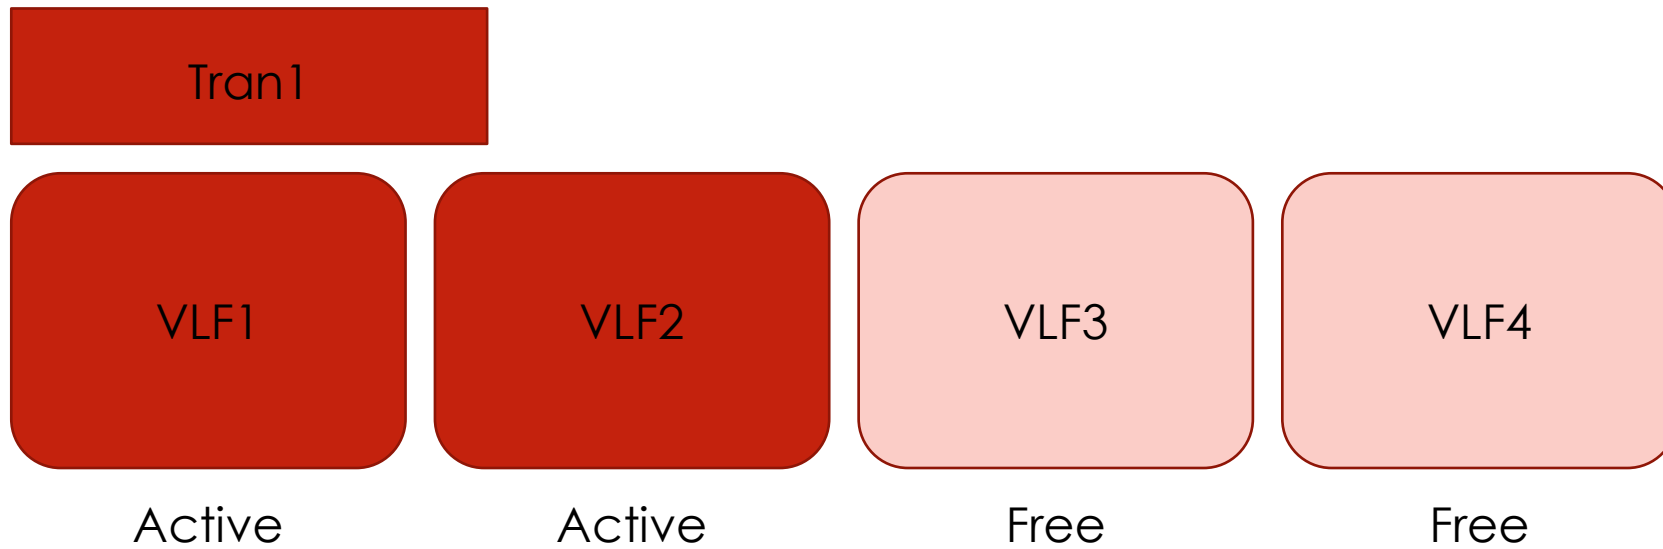
In FULL recovery, VLFs freed by LOG backup

Log records are also needed for rollback and high availability

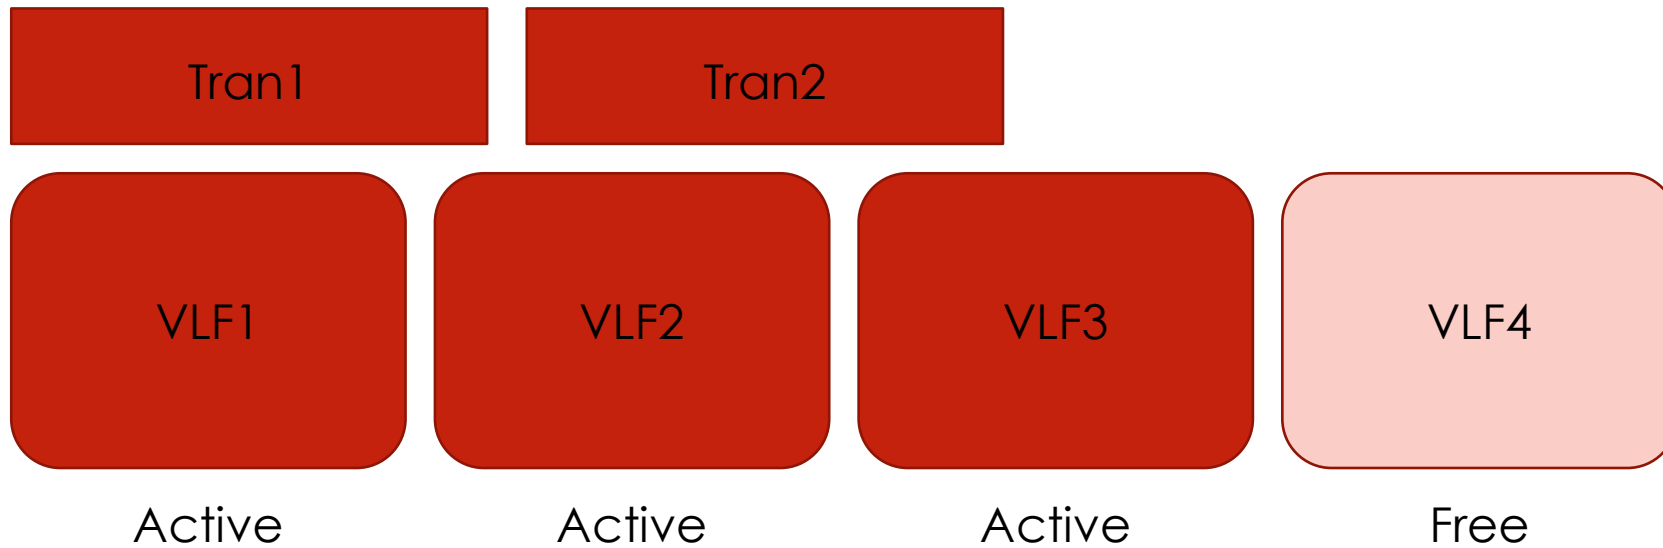
VIRTUAL LOG FILES



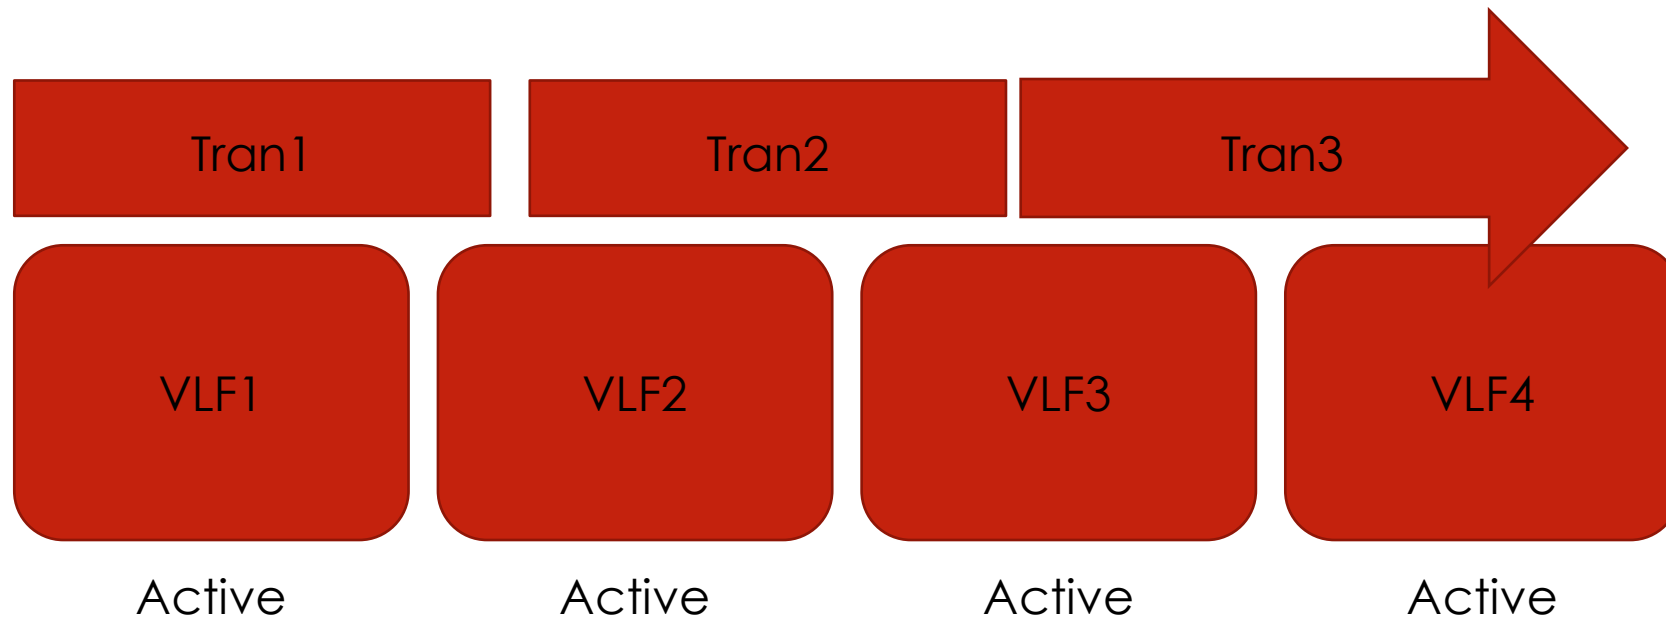
TRANSACTION 1 COMMITS



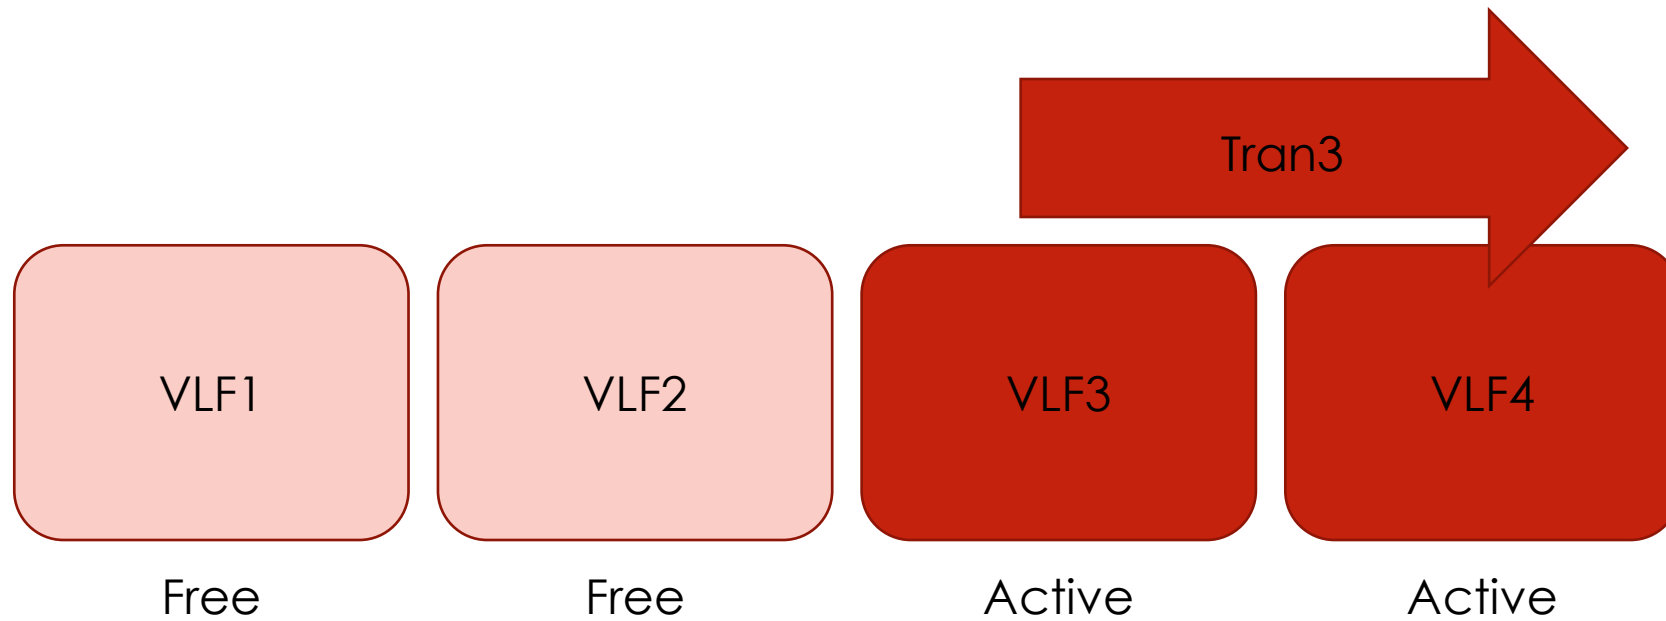
TRANSACTION 2 COMMITS



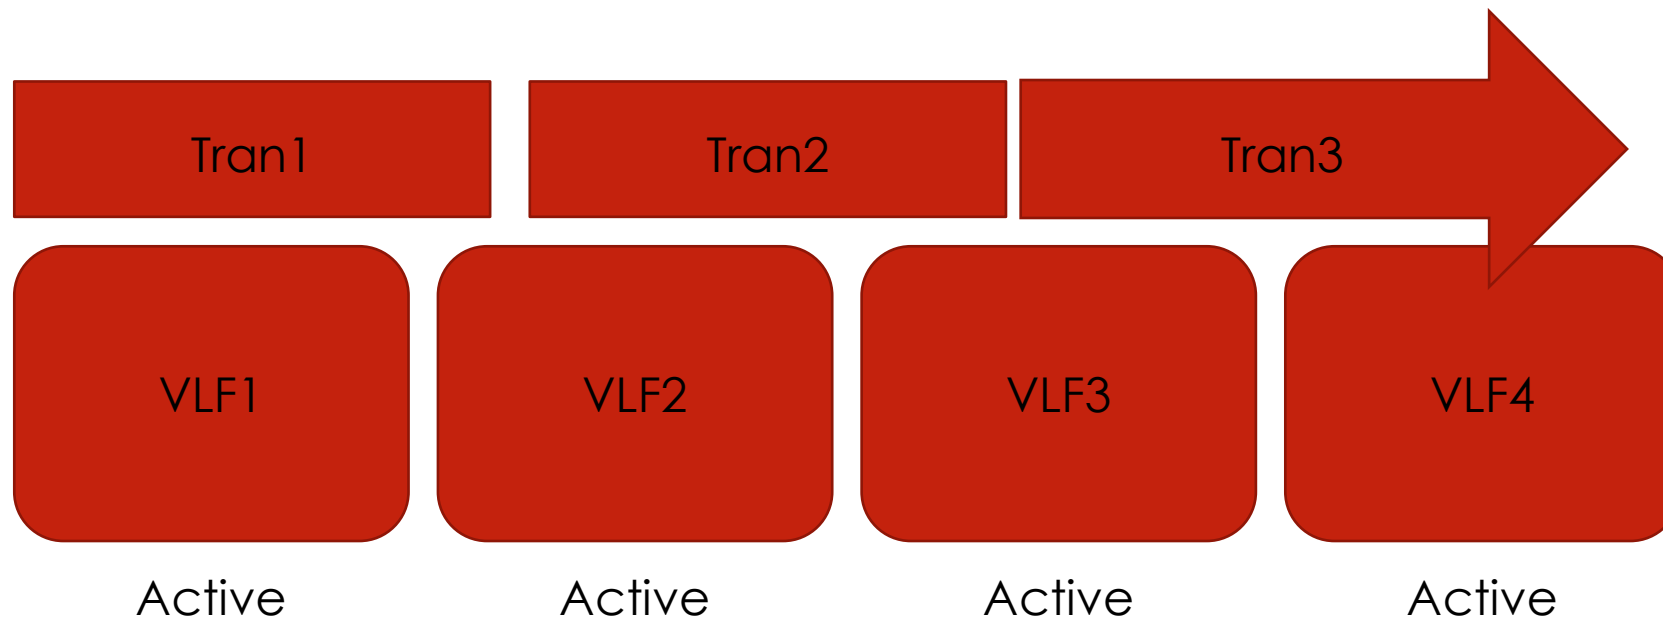
TRANSACTION 3 BEGINS




LOG BACKUP OR CHECKPOINT




NO LOG BACKUP OR CHECKPOINT



WITHOUT AUTOGROW

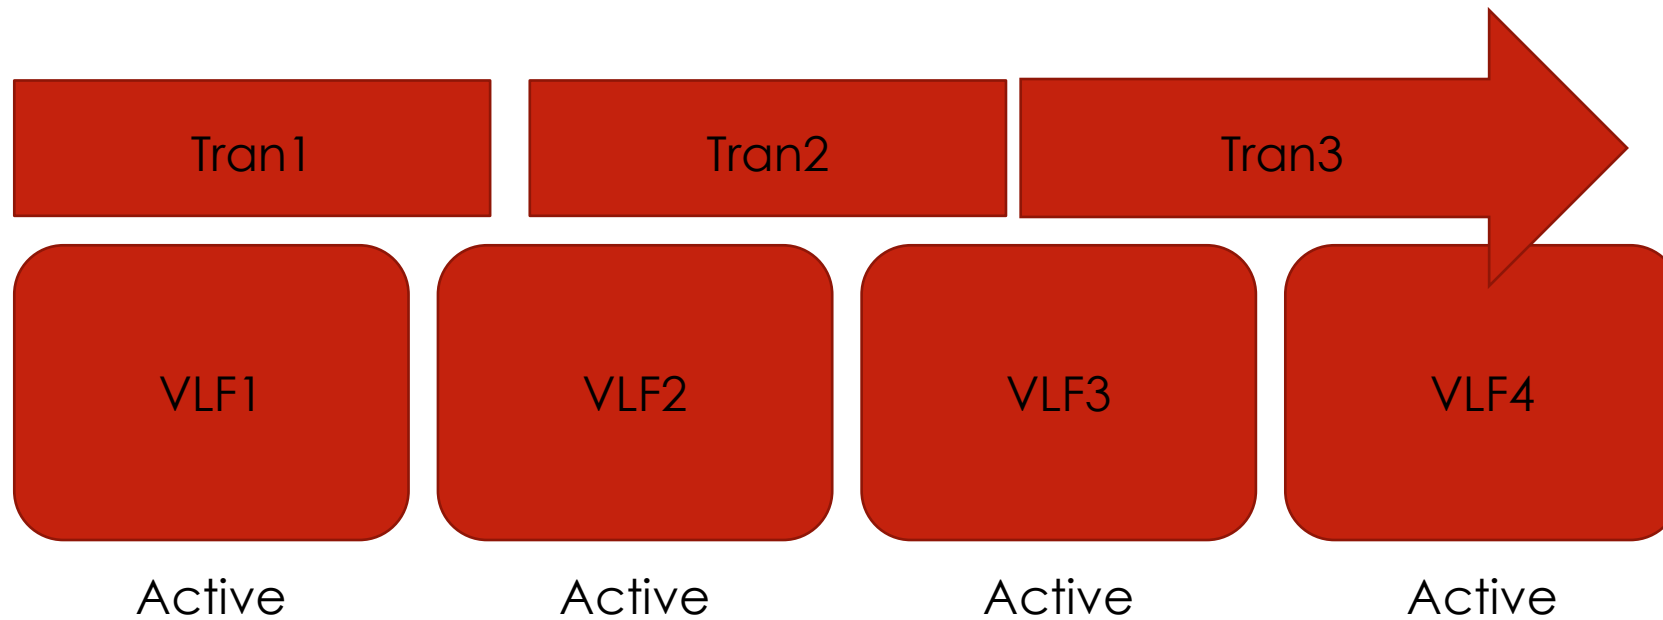




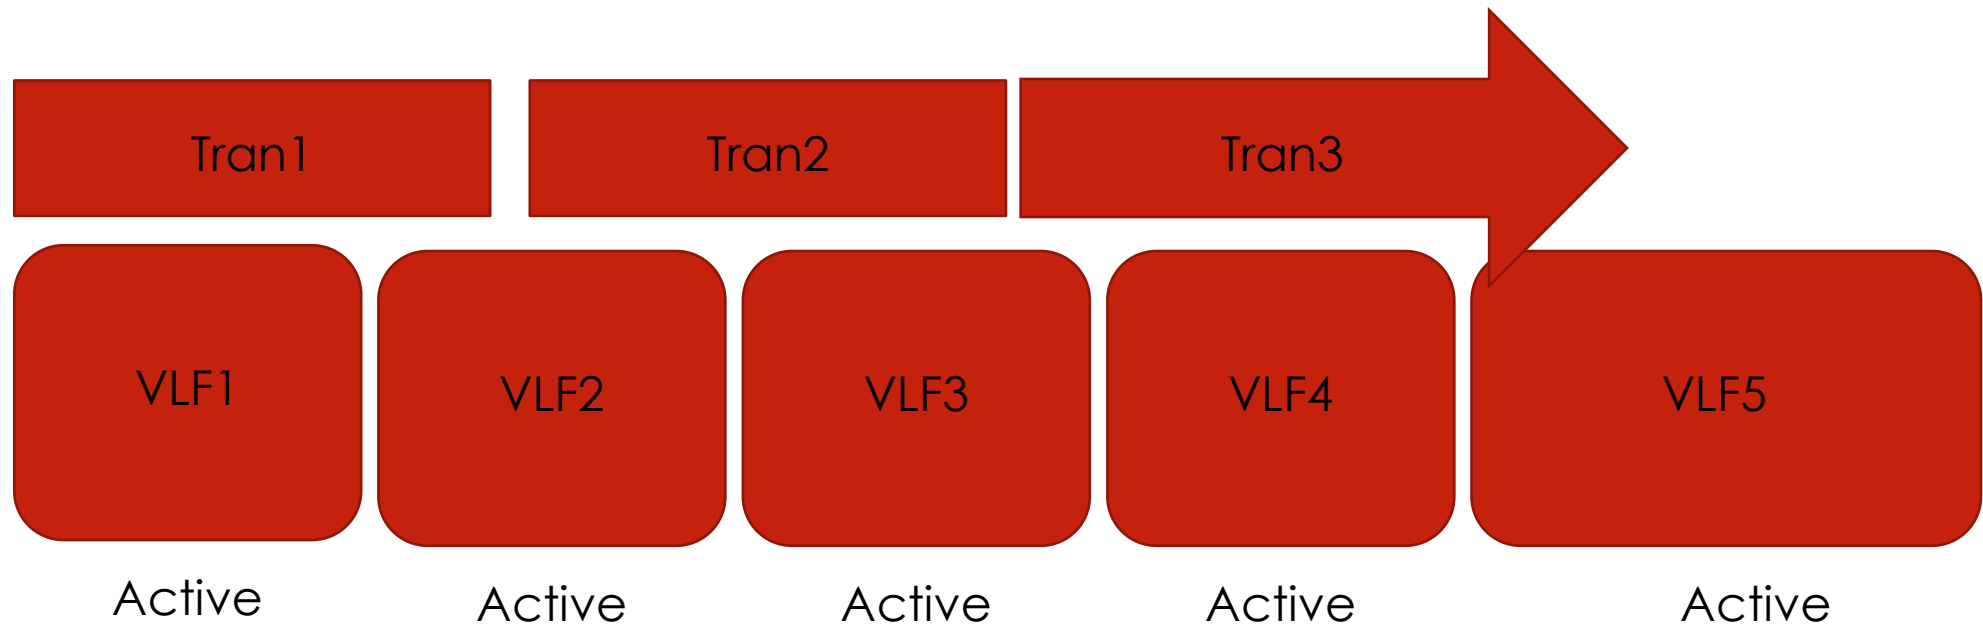
COMPUTER MALFUNCTION

Image Source: <https://tarnmoor.files.wordpress.com/2018/12/PICComputerMalfunction2.jpg>

WITH AUTOGROW



AUTOGROW ADDS VLF(S)





CRASH RECOVERY

SQL Server uses the transaction log to maintain consistency and durability
After the SQL Server service restarts, each database transaction log is scanned

CURRENT CRASH RECOVERY PROCESS

3 phases

Analysis – Scans log from last checkpoint searching for
Transactions written to log file but not to data file
Transactions not committed

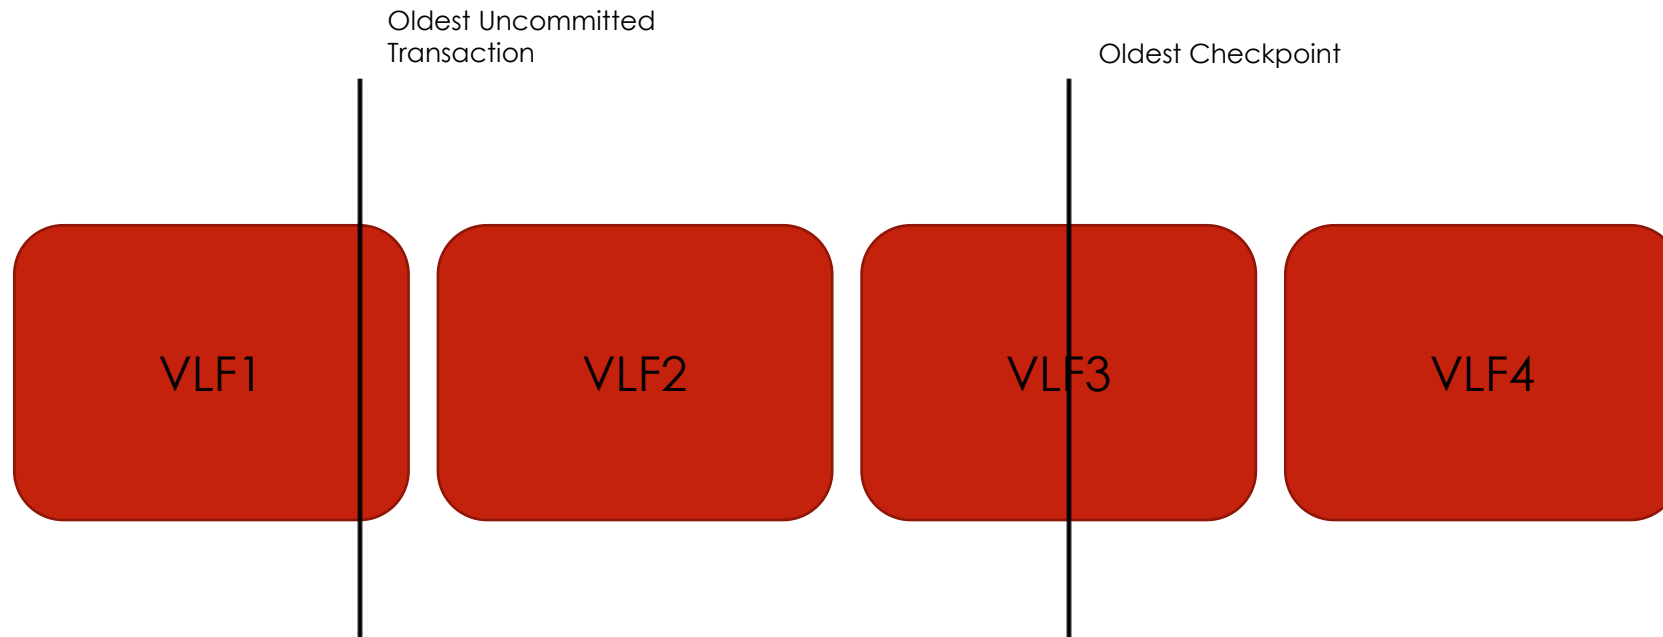
Redo

Committed transactions hardened to disk in data file

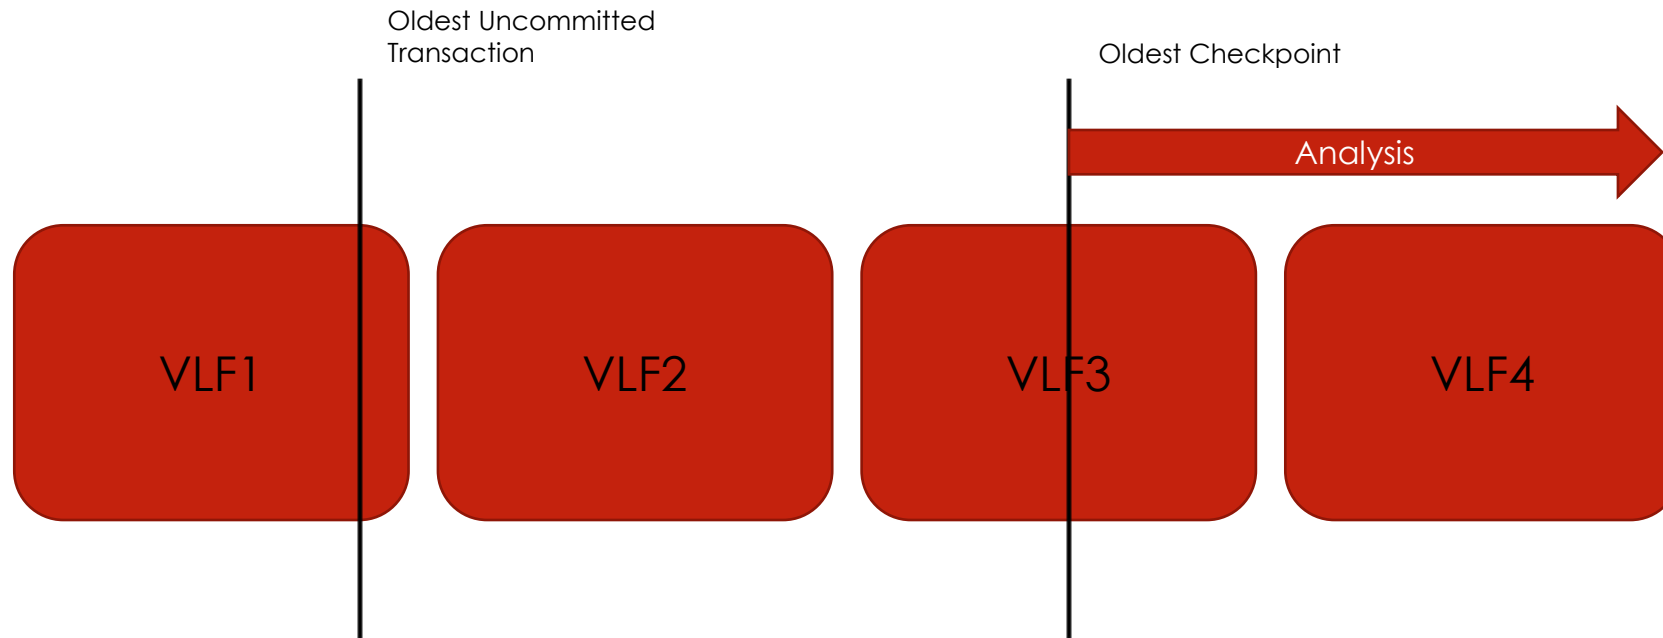
Undo

Uncommitted transactions rolled back

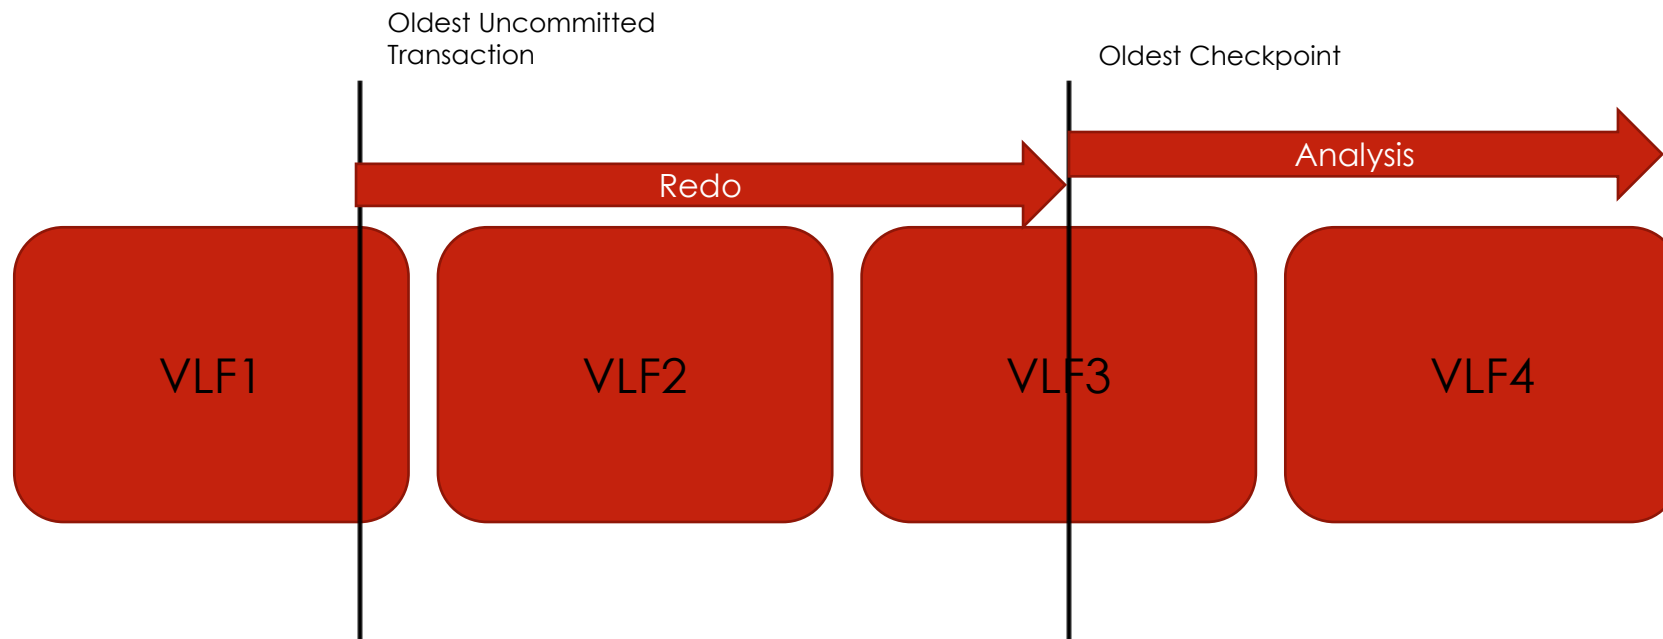
CRASH RECOVERY



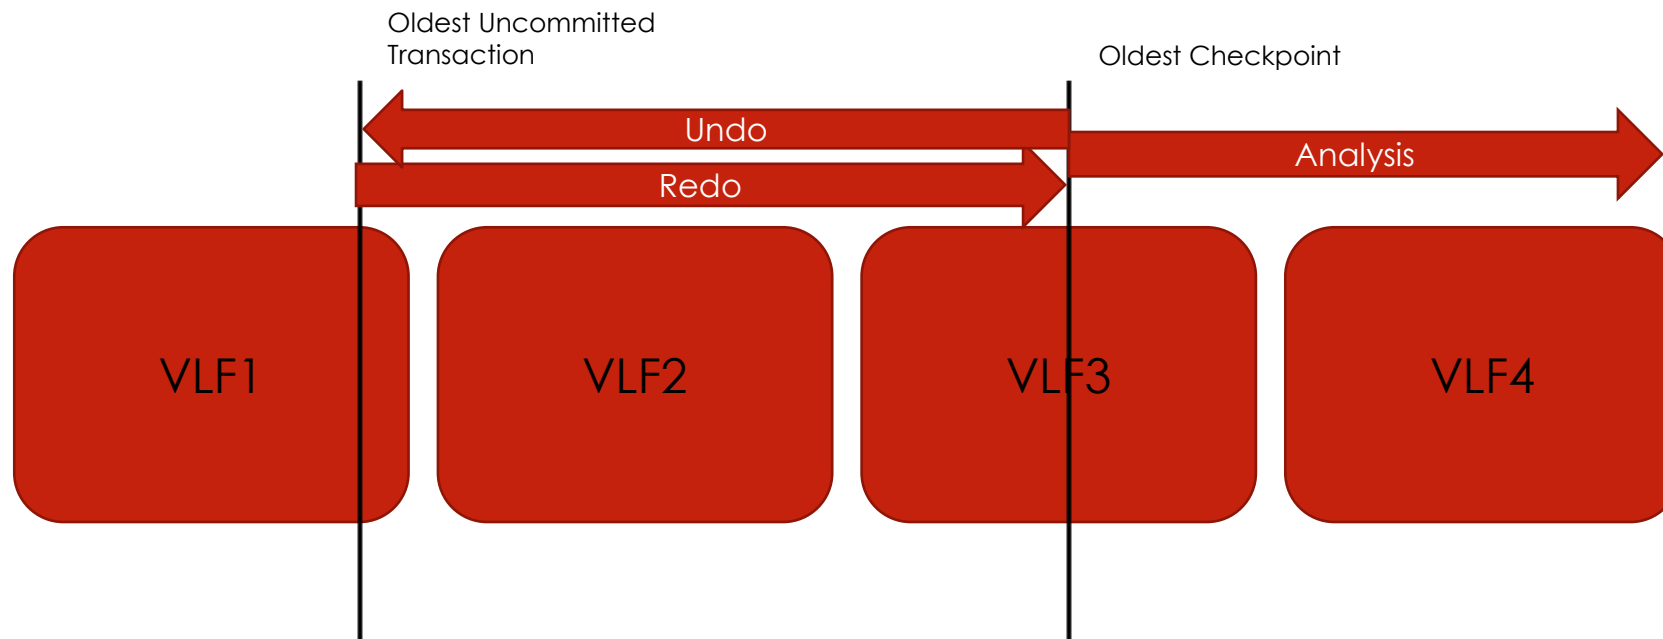
ANALYSIS



REDO



UNDO





ACCELERATED DATABASE RECOVERY – NEW CONCEPTS

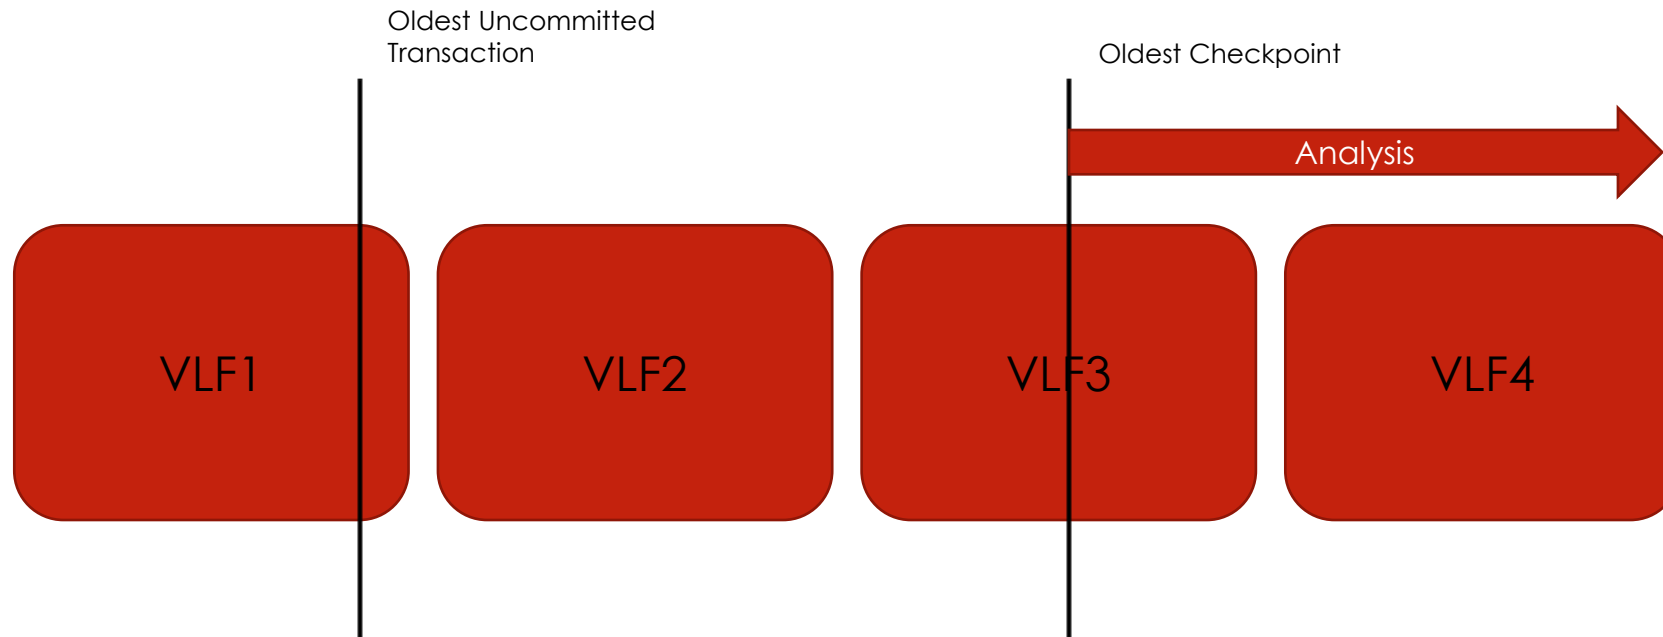
Persisted Version Store (PVS) – contains previous versions of modified rows, stored in the user database

Logical revert – On rollback, running transactions pull row version from the PVS

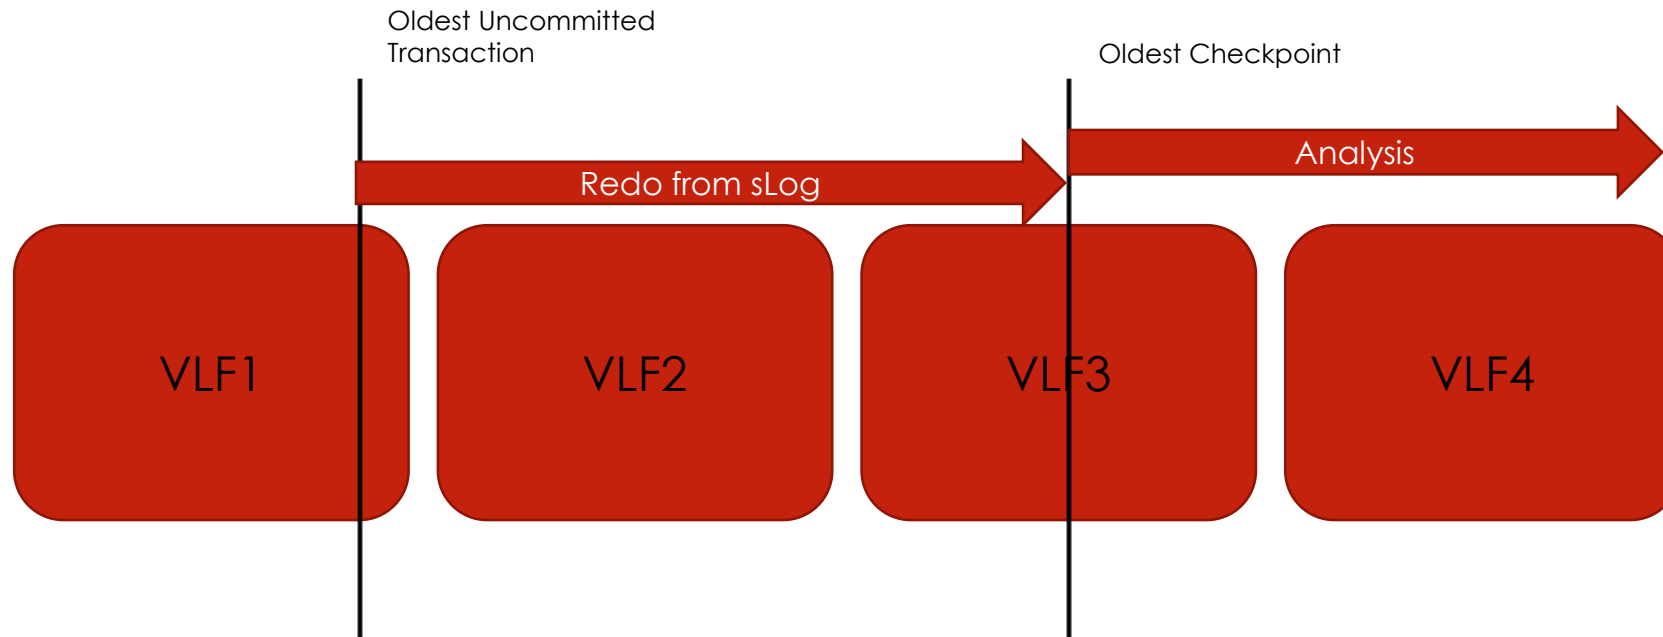
sLog – In-memory log stream that stores non-versioned activity (system metadata changes, locks for DDL, cache invalidation)

Cleaner – periodic process that cleans up unneeded row versions

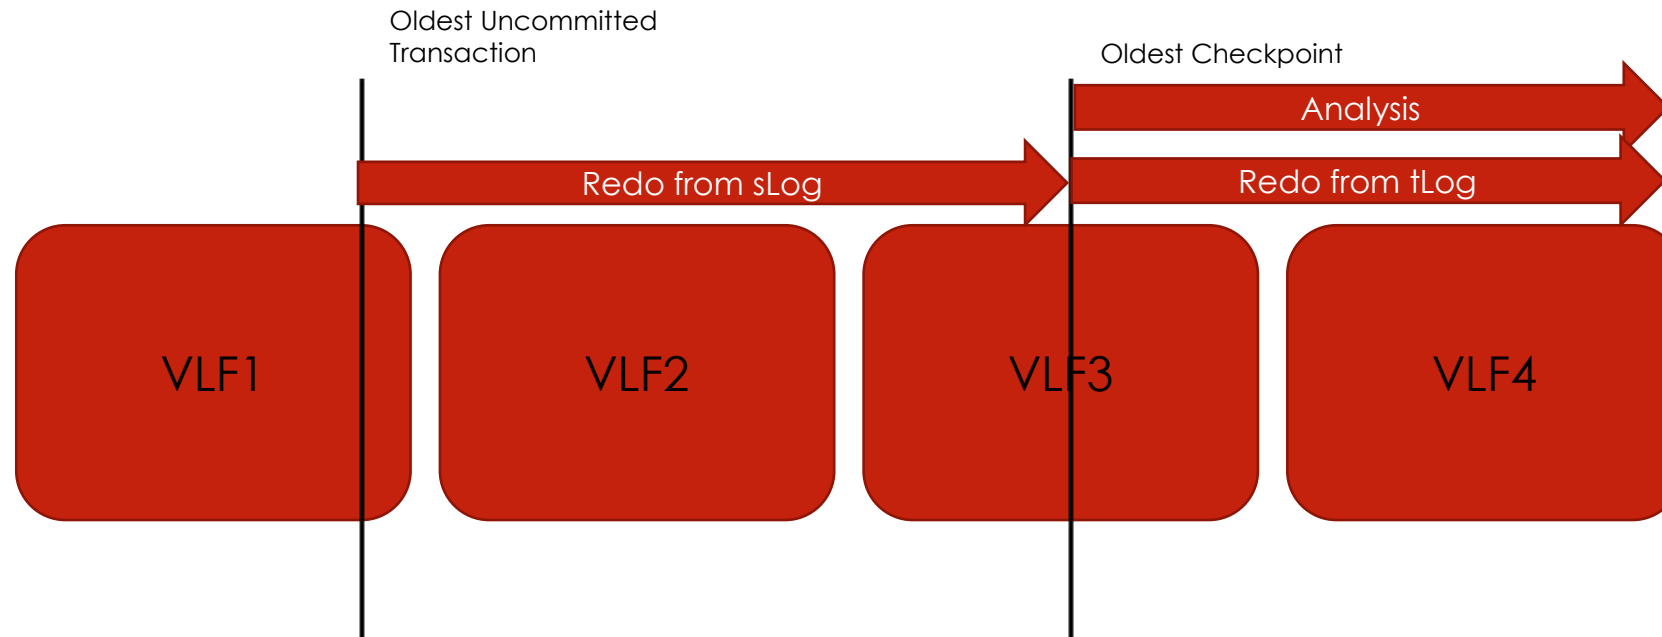
ADR ANALYSIS



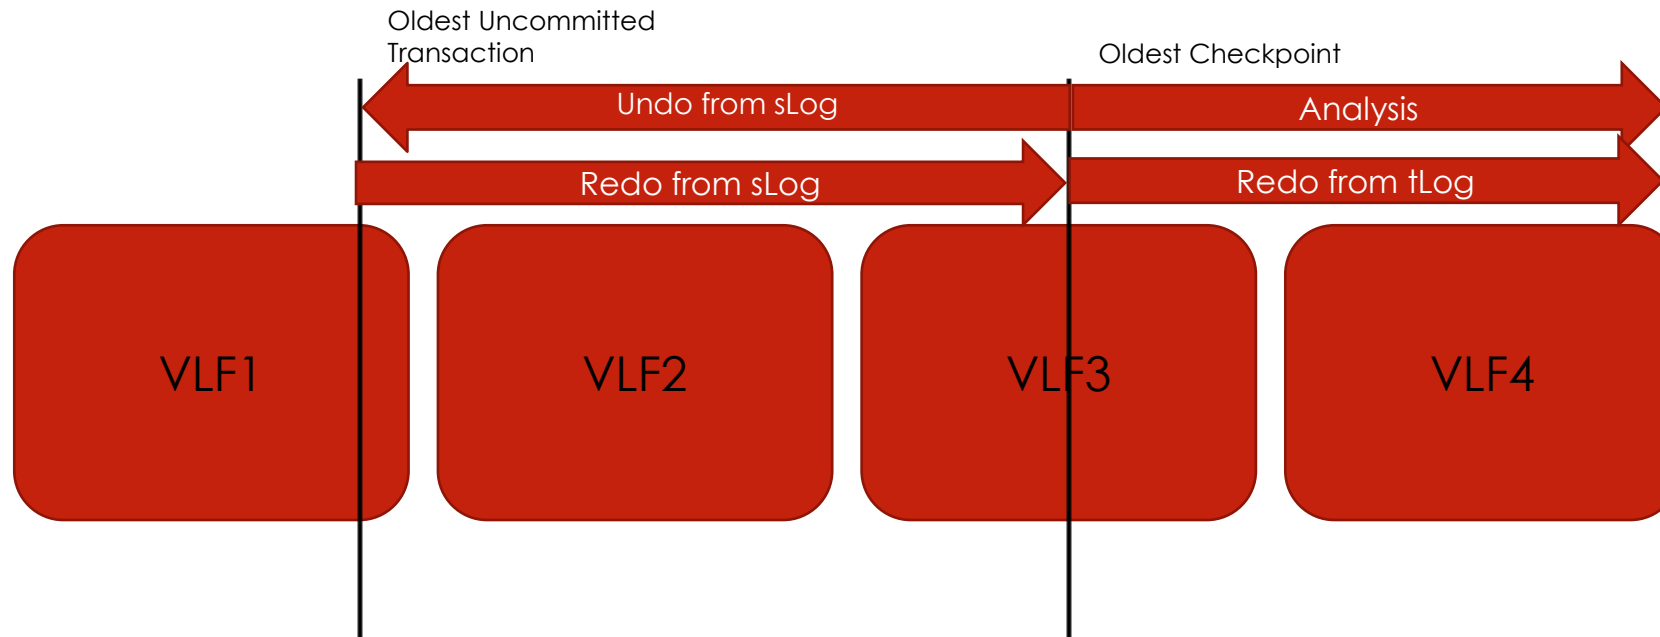
REDO FROM SLOG



REDO FROM TLOG



REDO FROM TLOG





BENEFITS OF ADR

Faster crash recovery

Faster AG failover

Faster rollback

Fast log truncation means smaller logs



BENEFITS OF ADR

Faster crash recovery

Faster AG failover

Faster rollback

Fast log truncation means smaller logs

AND!



BENEFITS OF ADR

Faster crash recovery

Faster AG failover

Faster rollback

Fast log truncation means smaller logs

AND!

It is available in Standard Edition!!!



Source: <http://www.quickmeme.com/meme/3q7ogg>



ACCELERATED DATABASE RECOVERY DEMO



IN-MEMORY TEMPDB METADATA TABLES



WHAT IS TEMPDB?

tempdb is one of the system databases in SQL Server

Used to create temporary objects

Used for sort space

Used by everybody

Can cause contention



ONE MAN'S OPINION

“TempDB – or as I call it.... SQL Server’s public toilet. You have no idea what other filthy, disgusting things people are doing in TempDB.” –Brent Ozar

Source: <https://ozar.me/2013/02/why-not-everybody-loves-my-sessions/>



THE OLD CONTENTION PROBLEM

Each data file contains pages to manage page allocation in the database
Because everyone uses tempdb, these pages can act as a bottleneck



THE OLD CONTENTION SOLUTION

Add data files to tempdb

This increases the number of metadata pages

SQL Server can parallelize activity



THE NEW CONTENTION PROBLEM

Each database contains system tables that store metadata about objects

Normally, this isn't a problem because objects are usually static

Because temp tables are constantly created and deleted, these tables can be bottlenecks

The problem is PAGELATCH waits



THE NEW CONTENTION SOLUTION

Create the tempdb system tables in memory

Eliminates latch contention



IN-MEMORY TEMPDB METADATA DEMO



HOW LATCHING WORKS

PAGE IS READ FROM DISK

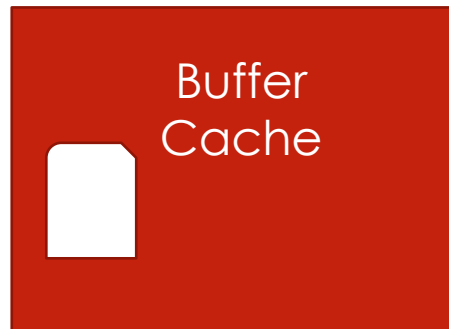
SQL Server
executes an
UPDATE to a
page not in the
buffer cache

Buffer
Cache



DATA IS READ

Page read into
the buffer from
disk



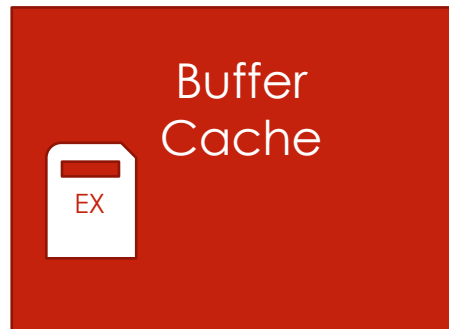
PAGELATCH_EX IS TAKEN

SQL Server takes an exclusive latch on the page to prevent collisions in-memory
While the exclusive latch is held no other processes can access the page
In the demo, all activity is updating the same row, so a single page is a bottleneck



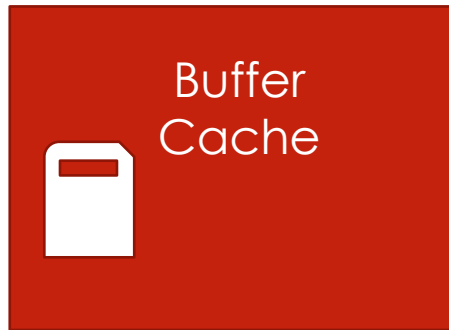
DATA IS WRITTEN TO THE PAGE

SQL Server
updates data on
the page



LATCH IS RELEASED

Latch is released
and update
completes





TEMPDB LATCHING CONTENTION

tempdb metadata pages are accessed for each creation or deletion of a temp table

Each access requires a latch

Latch contention can occur

Negative impact to performance

IN-MEMORY ROW STRUCTURE

Row structure

Row header

Row data

ROW HEADER

Row structure

Row header

Inserted Row Data



Start TS

End TS

StatementID

idxLinkCount

of Indexes

IN-MEMORY INSERT

Row structure

Row header

Inserted Row Data



50

∞

1

idxLinkCount

of Indexes

IN-MEMORY UPDATE

Row structure

Row header

Updated Row Data



50

90

1

idxLinkCount

of Indexes

90

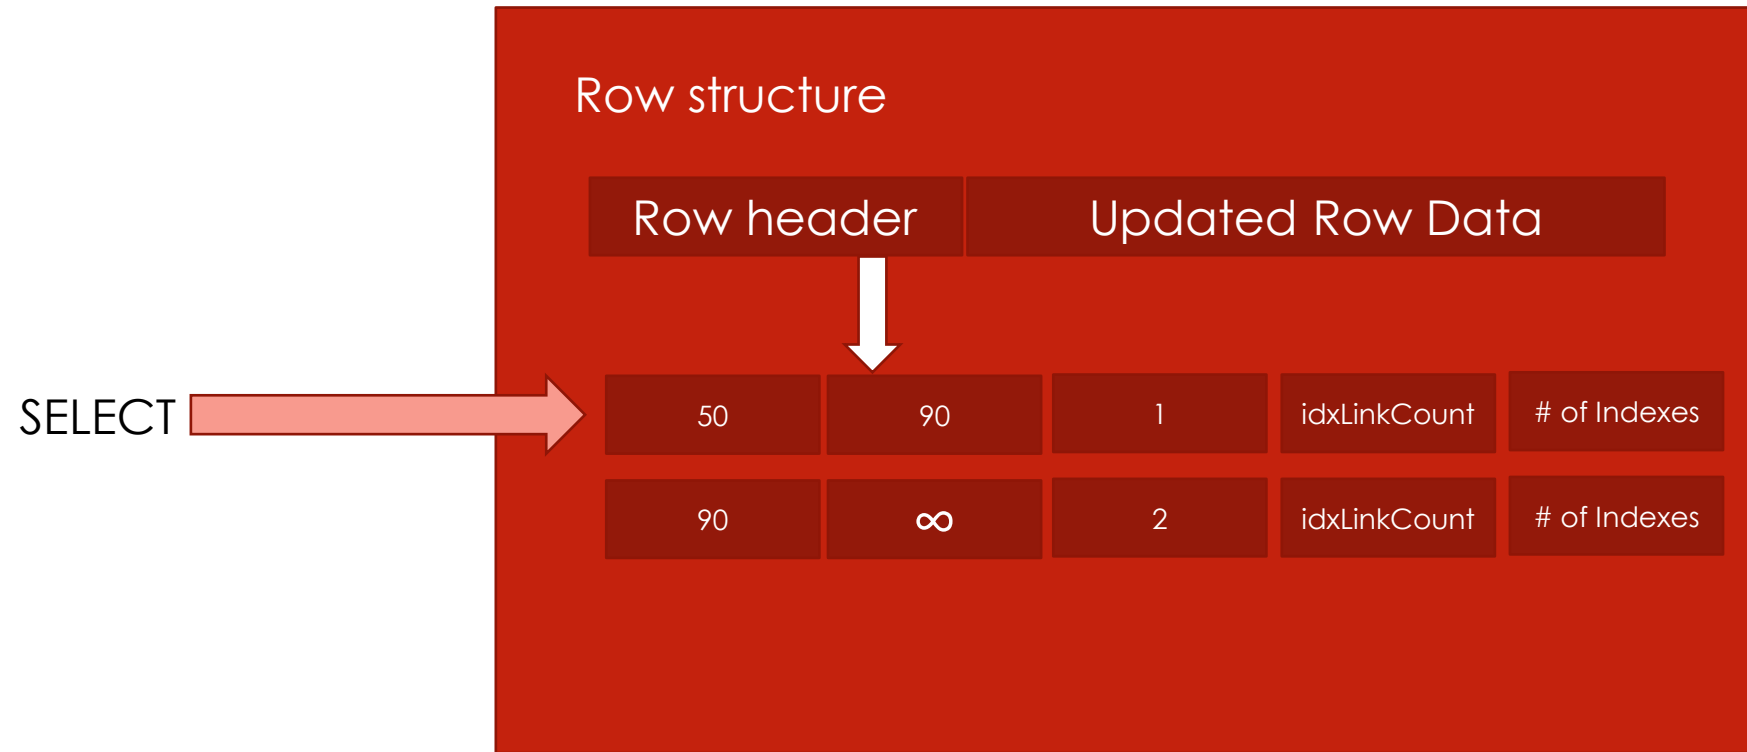
∞

2

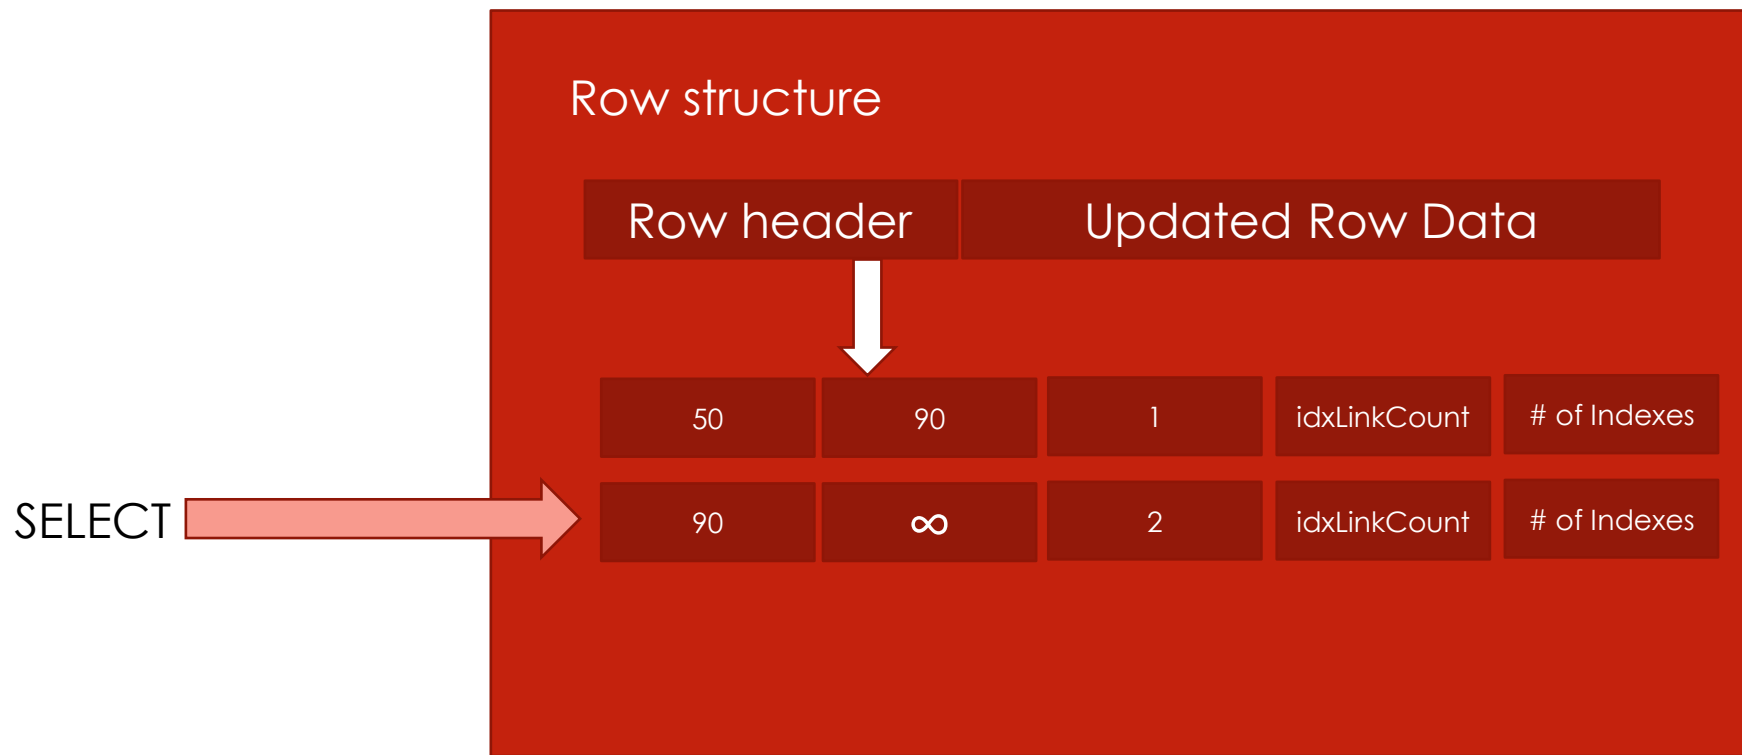
idxLinkCount

of Indexes

SELECT WITH TIMESTAMP 60



SELECT WITH TIMESTAMP 120





PERSISTENT MEMORY



BEFORE PERSISTENT MEMORY

Historically, RAM has been transient

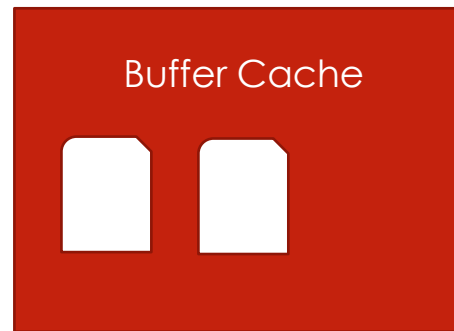
On shutdown, data in RAM is lost

As a result, the transaction log buffer is flushed to disk on COMMIT

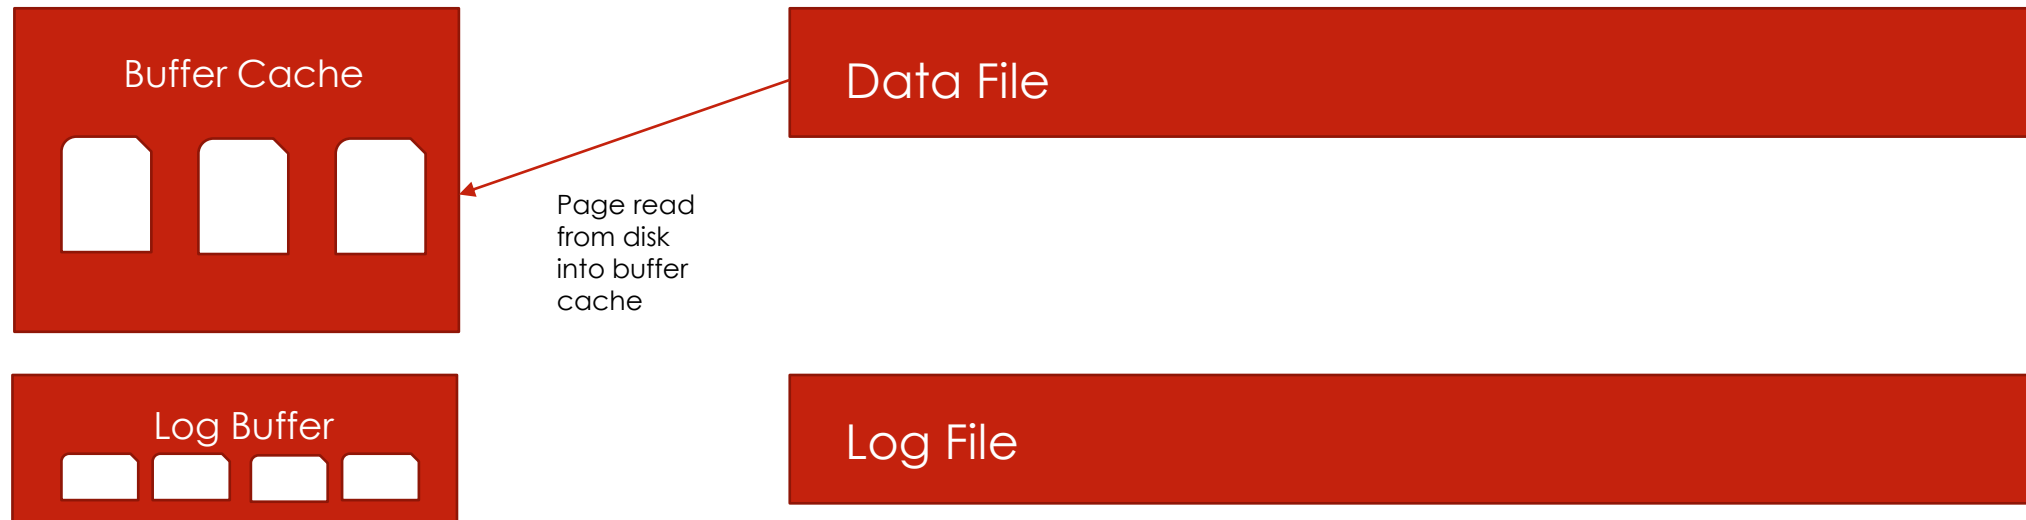
SQL Server must wait for confirmation that the flush has completed

This allows redo and undo to take place in the event of a crash


LOG BUFFER AND TRANSACTIONS



READ PAGE INTO BUFFER CACHE



CHANGE PAGE IN MEMORY



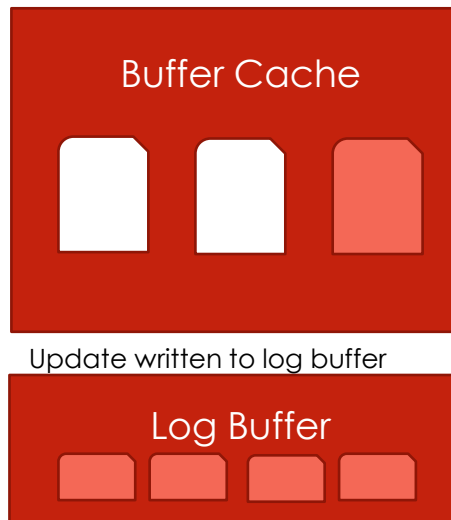
Page updated



Data File

Log File

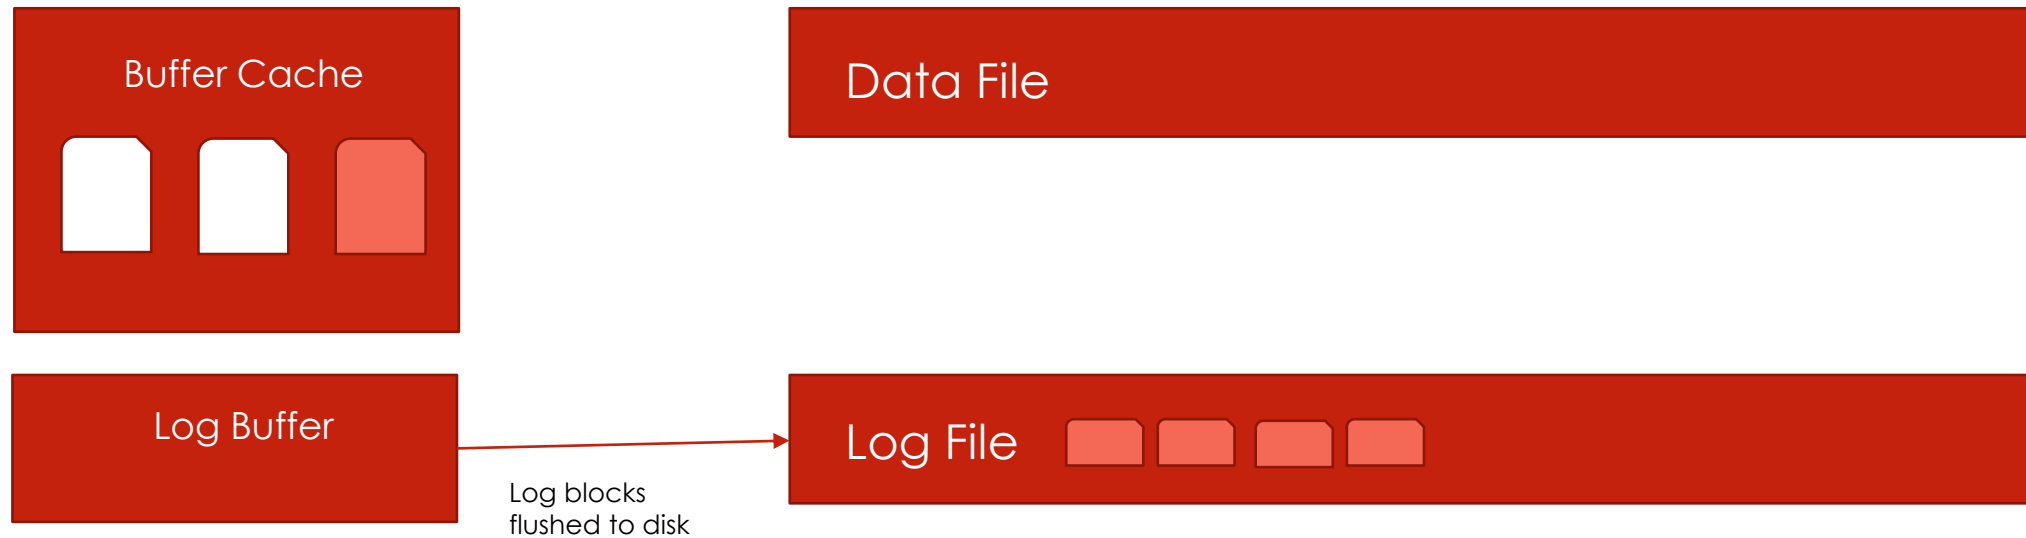
WRITE LOG RECORDS TO LOG CACHE



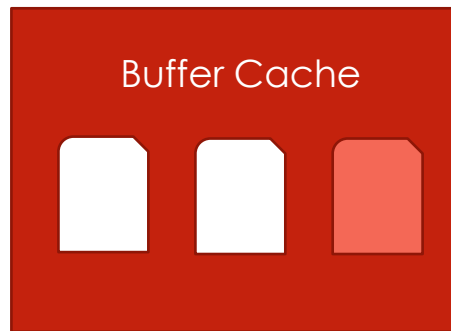
Data File

Log File

FLUSH LOG CACHE TO DISK



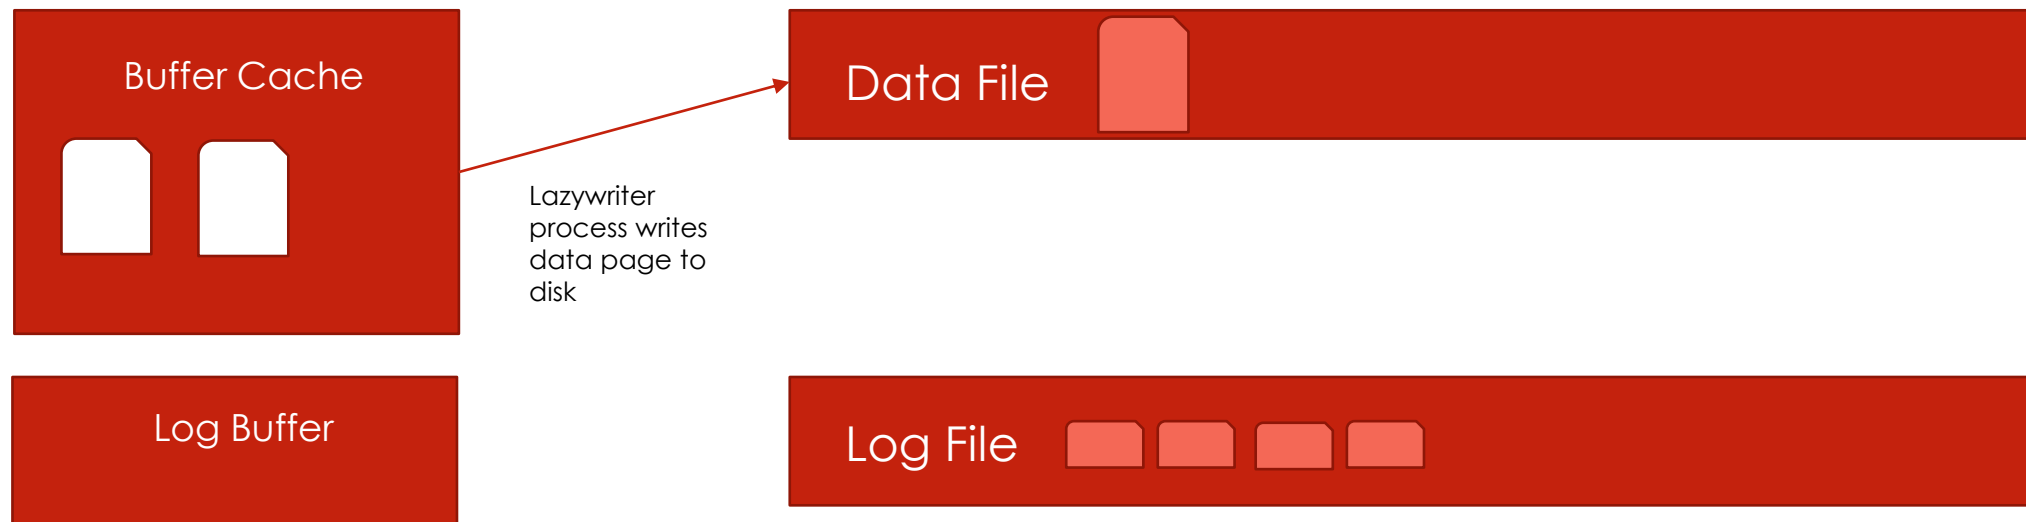
WRITE LOG RECORDS TO LOG CACHE



Transaction commits



DATA PAGES FLUSHED ASYNCHRONOUSLY





PERSISTENT MEMORY

New development in hardware

RAM with a battery

Data stored in RAM can survive a restart

This provides several opportunities for performance enhancements




PERSISTENT LOG BUFFER CACHE

Log flush is no longer necessary on COMMIT

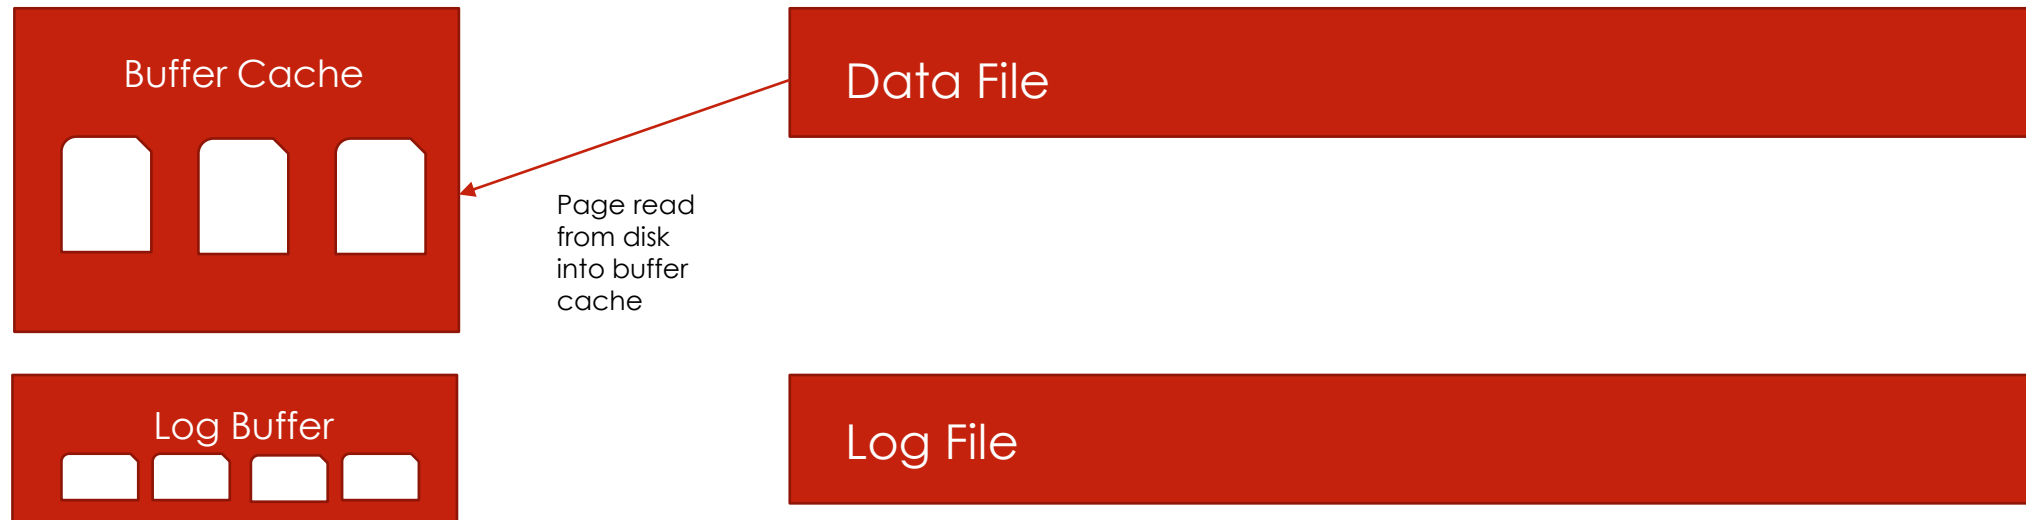
Log flush can happen in the background

Persistent buffer processed with the log file on restart

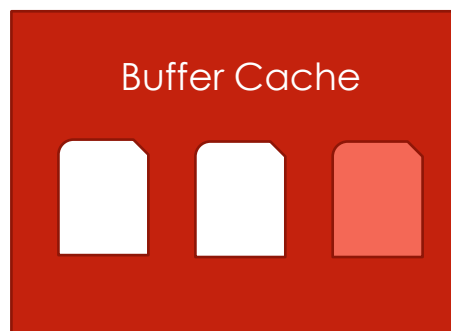
PERSISTENT LOG BUFFER CACHE



READ PAGE INTO BUFFER CACHE



CHANGE PAGE IN MEMORY



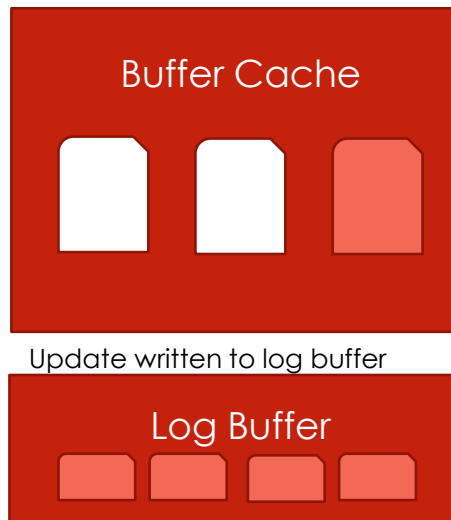
Page updated



Data File

Log File


WRITE LOG RECORDS TO LOG CACHE




Data File

Log File

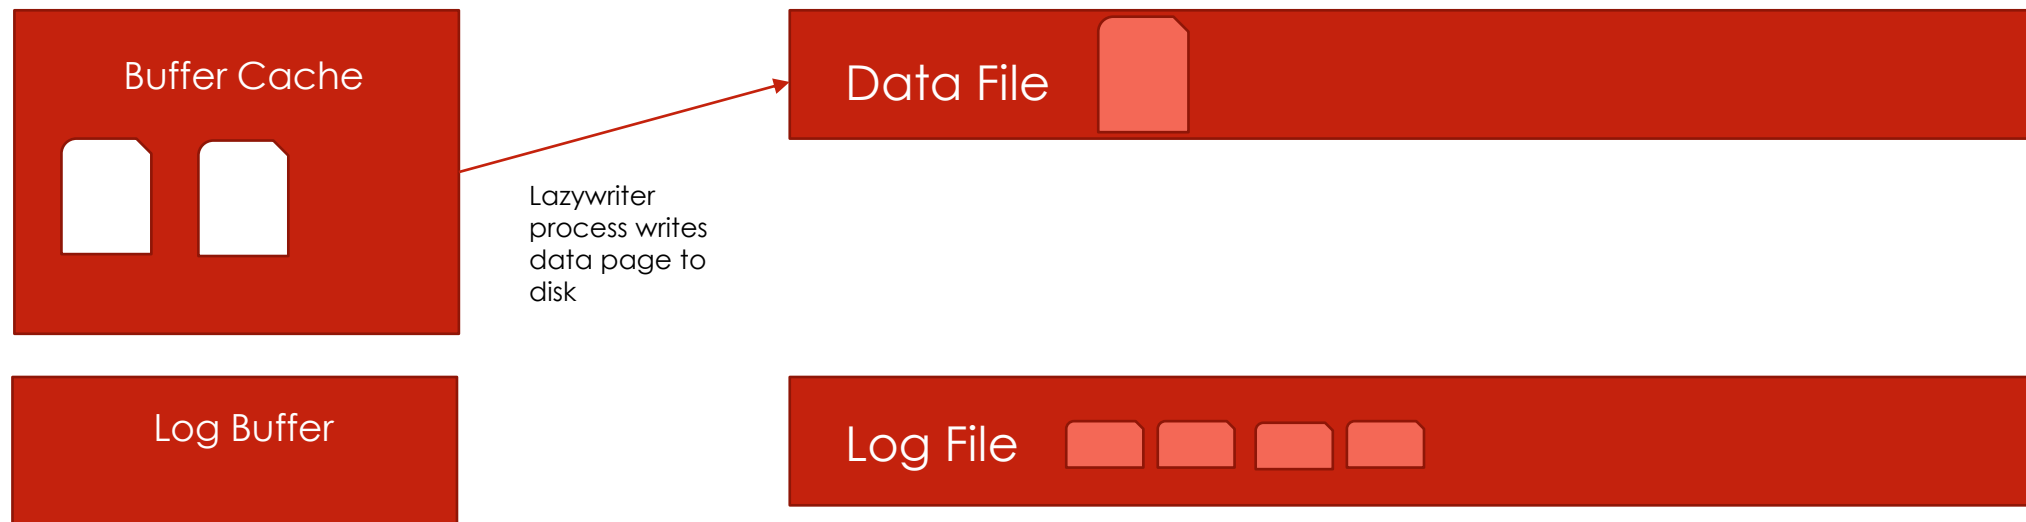
TRANSACTION COMMITS



LOG CACHE FLUSHED BY BACKGROUND PROCESS



DATA PAGES FLUSHED ASYNCHRONOUSLY





ADVANTAGES OF PERSISTENT LOG BUFFER CACHE

Log buffers not flushed to disk on commit

Speeds transactional processing

Reduces LOGWRITE waits



HYBRID BUFFER POOL

Enhancement of Buffer Pool Extension (BPE)

BPE extended the buffer pool onto fast SSD disk

Hybrid Buffer Pool extends the buffer pool onto PMEM

BPE without the IO overhead



ENLIGHTENED IO

Linux-only feature

Data and log files can be placed on PMEM

Allows the file system and storage stack to be bypassed

Note – PMEM is not as fast as traditional RAM

If your database fits in memory, this is not the solution for you




NO PERSISTENT MEMORY DEMO

WHY?

[Home](#) > [Data Storage Products](#) > [Hard Drives](#) > [Solid State Drives \(SSDs\)](#)



 Share

 Save as Favorites



Intel Optane DC Persistent – DDR-T – module – 512 GB – DIMM 288-pin

Mfg.Part: UCS-MP-512GS-A0= | CDW Part: 6020774 | UNSPSC: 43201402

Availability: • 4-6 Weeks

Orders placed today will ship within 4-6 weeks by a CDW partner.

 Advertised Price

Claim up to a 5% Discount
[Create an account](#) to get My CDW Advantage today.

1

Add to Cart

[Lease Option](#) (\$562.19/month)

Product Details

- DDR-T
- module
- 512 GB
- DIMM 288-pin
- 2666 MHz / PC4-21300
- 1.2 V
- for UCS C220 M5

[View Full Product Details](#)

Source: [Intel Optane DC Persistent - DDR-T - module - 512 GB - DIMM 288-pin - UCS-MP-512GS-A0= - Hard Drives - CDW.com](#)

THAT'S WHY

[Home](#) > [Data Storage Products](#) > [Hard Drives](#) > [Solid State Drives \(SSDs\)](#)



Share



Save as Favorites



Intel Optane DC Persistent – DDR-T – module – 512 GB – DIMM 288-pin

Mfg.Part: UCS-MP-512GS-A0= | CDW Part: 6020774 | UNSPSC: 43201402

Availability: • **4-6 Weeks**

Orders placed today will ship within 4-6 weeks by a CDW partner.

\$19,206.99 Advised Price

Claim up to a 5% Discount

[Create an account](#) to get My CDW Advantage today.

1



Add to Cart

[Lease Option](#) (\$562.19/month)

Product Details

- DDR-T
- module
- 512 GB
- DIMM 288-pin
- 2666 MHz / PC4-21300
- 1.2 V
- for UCS C220 M5

[View Full Product Details](#)



WHAT WE'VE COVERED

Accelerated Database Recovery

In-memory tempdb metadata

Persistent Memory

- Persistent Log Buffer Cache

- Hybrid Buffer Pool

- Enlightened IO

RESOURCES

Notebook Files – <https://github.com/skreebydba/MinnesotaPresentation>

Tiger Team Materials – <https://microsoft.github.io/sqlworkshops/>

Hybrid Buffer Pool –

<https://docs.microsoft.com/en-us/sql/database-engine/configure-windows/hybrid-buffer-pool?view=sql-server-2017>

RESOURCES

Accelerated Database Recovery –

<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-accelerated-database-recovery>

Constant Time Recovery in Azure SQL Database (White paper that gets deep into the internals of ADR, also known as CTR) –

<https://www.microsoft.com/en-us/research/publication/constant-time-recovery-in-azure-sql-database/>

In-memory tempdb Metadata –

<https://docs.microsoft.com/en-us/sql/relational-databases/databases/tempdb-database?view=sql-server-ver15#memory-optimized-tempdb-metadata>